

```
from google.colab import drive
drive.mount('/gdrive/')
%cd /gdrive
```

```
Mounted at /gdrive/
/gdrive
```

```
ls
```

```
MyDrive/ Shareddrives/
```

```
cd/gdrive/My Drive/Multinomial animal classification/
```

```
/gdrive/My Drive/Multinomial animal classification
```

```
ls
```

```
cats/ dogs/ horses/
```

```
import tensorflow as tf
from tensorflow.keras import models, layers
import matplotlib.pyplot as plt
from PIL import Image
import numpy as np
from tensorflow.keras.layers.experimental import preprocessing
from tensorflow.keras.preprocessing.image import load_img
```

```
import os
import PIL
import pathlib
import pandas as pd
import numpy as np
import matplotlib.image as mpimg
import matplotlib.pyplot as plt
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import preprocessing
from tensorflow.keras import layers
from tensorflow.keras.models import Sequential
from keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPool2D
from tensorflow.keras.optimizers import RMSprop
from tensorflow.keras.models import Sequential
from tensorflow.keras.layers import Dense, Dropout, Flatten, Conv2D, MaxPooling2D
```

```

from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.python.ops.numpy_ops import np_utils

#from google.colab import drive
#drive.mount('/content/drive')

```

```

BATCH_SIZE = 62
IMAGE_SIZE = 256
EPOCHS=50
CHANNELS=3

```

```

dataset = tf.keras.preprocessing.image_dataset_from_directory(
    "/gdrive/My Drive/Multinomial animal classification/",
    seed=123,
    shuffle=True,
    image_size=(IMAGE_SIZE, IMAGE_SIZE),
    batch_size=BATCH_SIZE
)

```

Found 606 files belonging to 3 classes.

```

class_names = dataset.class_names
class_names

```

```
['cats', 'dogs', 'horses']
```

```
len(dataset)
```

```
10
```

```

for image_batch, label_batch in dataset.take(1):
    print(image_batch.shape)
    print(image_batch[1])
    print(label_batch.numpy())

```

```

(62, 256, 256, 3)
tf.Tensor(
[[[ 17.091797  13.091797  28.091797]
   [ 33.476562  30.214844  44.845703]
   [ 39.691406  37.691406  51.691406]
   ...
   [188.69336  161.07812  140.23242 ]
   [190.47656  166.63086  146.36914 ]
   [186.36914  163.49219  146.87695 ]]]

[[ 51.65497  47.65497  62.65497 ]
 [ 56.120422  52.858704  67.48956 ]
 [ 45.33954  43.33954  57.33954 ]

```

```

...
[189.45978 161.84454 140.99884 ]
[188.9986 165.1529 144.89117 ]
[185.35138 162.47443 144.26971 ]]

[[ 73.170044 69.170044 84.170044]
 [ 65.42218 61.537537 76.168396]
 [ 36.514038 32.826538 46.826538]
...
[195.31512 167.69989 146.85419 ]
[193.06445 169.21875 148.95703 ]
[190.20477 167.32782 148.95868 ]]

...

[[135.9151 146.07135 148.07135 ]
 [145.86835 155.344 157.65546 ]
 [142.30444 150.42468 153.36456 ]
...
 [ 39.55902 37.71527 23.934021]
 [ 37.1474 36.900513 21.5224 ]
 [ 46.068054 46.755554 30.443054]]

[[128.05878 139.05878 141.05878 ]
 [141.04938 151.68024 153.68024 ]
 [141.09961 149.9845 152.54205 ]
...
 [ 39.854553 38.739197 20.724731]
 [ 37.339294 36.339294 18.151794]
 [ 45.71466 44.71466 26.52716 ]]

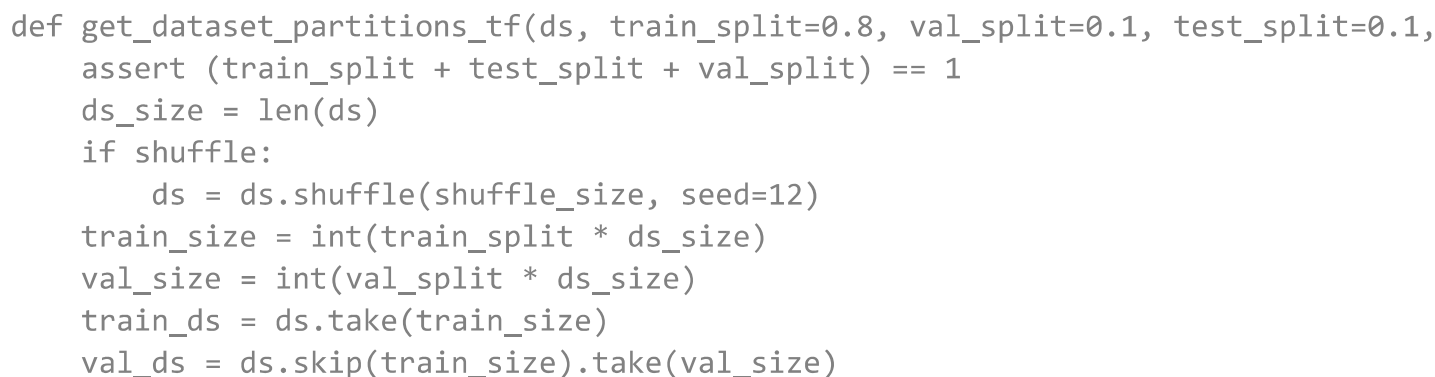
[[114.353516 125.353516 127.353516]
 [126.21484 136.8457 138.8457 ]
 [125.69336 135.69336 137.69336 ]
...
 [ 29.53711 27.30664 9.152344]
 [ 25.892578 24.892578 4.892578]
 [ 33.13867 32.13867 12.138672]]], shape=(256, 256, 3), dtype=float32)
[1 1 1 0 1 1 1 1 0 2 1 0 0 2 0 1 1 1 0 0 2 0 1 2 2 1 0 2 0 2 2 1 2 2 2 2 1
 0 0 0 1 0 1 2 1 0 2 2 0 2 2 2 1 0 1 2 0 2 2 2 1 1]

```

```

plt.figure(figsize=(15, 15))
for image_batch, labels_batch in dataset.take(1):
    for i in range(BATCH_SIZE):
        ax = plt.subplot(8, 8, i + 1)
        plt.imshow(image_batch[i].numpy().astype("uint8"))
        plt.title(class_names[labels_batch[i]])
        plt.axis("off")

```



```
test_ds = ds.skip(train_size).skip(val_size)
# Autotune all the 3 datasets
train_ds = train_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
val_ds = val_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
test_ds = test_ds.cache().shuffle(1000).prefetch(buffer_size=tf.data.AUTOTUNE)
return train_ds, val_ds, test_ds
```

```
train_ds, val_ds, test_ds = get_dataset_partitions_tf(dataset)
```

```
resize_and_rescale = tf.keras.Sequential([
    layers.experimental.preprocessing.Resizing(IMAGE_SIZE, IMAGE_SIZE),
    layers.experimental.preprocessing.Rescaling(1./255),
])
```

```
data_augmentation = tf.keras.Sequential([
    layers.experimental.preprocessing.RandomFlip("horizontal_and_vertical"),
    layers.experimental.preprocessing.RandomRotation(0.2),
])
```

```
input_shape = (BATCH_SIZE, IMAGE_SIZE, IMAGE_SIZE, CHANNELS)
n_classes = 9
```

```
model = models.Sequential([
    resize_and_rescale,
    # data_augmentation,
    layers.Conv2D(32, kernel_size = (3,3), activation='relu', input_shape=input_shape),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, kernel_size = (3,3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Conv2D(64, (3, 3), activation='relu'),
    layers.MaxPooling2D((2, 2)),
    layers.Flatten(),
    layers.Dense(64, activation='relu'),
    layers.Dense(n_classes, activation='softmax'),
])
model.build(input_shape=input_shape)
```

```
model.compile(
    optimizer='adam',
    loss=tf.keras.losses.SparseCategoricalCrossentropy(from_logits=False),
    metrics=['accuracy']
)
```

```
model.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape	Param #
=====		
sequential (Sequential)	(62, 256, 256, 3)	0
conv2d (Conv2D)	(62, 254, 254, 32)	896
max_pooling2d (MaxPooling2D)	(62, 127, 127, 32)	0
conv2d_1 (Conv2D)	(62, 125, 125, 64)	18496
max_pooling2d_1 (MaxPooling2D)	(62, 62, 62, 64)	0
conv2d_2 (Conv2D)	(62, 60, 60, 64)	36928
max_pooling2d_2 (MaxPooling2D)	(62, 30, 30, 64)	0
conv2d_3 (Conv2D)	(62, 28, 28, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(62, 14, 14, 64)	0
conv2d_4 (Conv2D)	(62, 12, 12, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(62, 6, 6, 64)	0
conv2d_5 (Conv2D)	(62, 4, 4, 64)	36928
max_pooling2d_5 (MaxPooling2D)	(62, 2, 2, 64)	0
flatten (Flatten)	(62, 256)	0
dense (Dense)	(62, 64)	16448
dense_1 (Dense)	(62, 9)	585

```
=====
Total params: 184,137
Trainable params: 184,137
Non-trainable params: 0
```

```
history = model.fit(  
    train_ds,  
    batch_size=BATCH_SIZE,  
    validation_data=val_ds,  
    verbose=1,  
    epochs=EPOCHS,  
)
```

```
Epoch 1/50  
8/8 [=====] - 49s 5s/step - loss: 1.7104 - accuracy: 0.3105  
Epoch 2/50  
8/8 [=====] - 37s 5s/step - loss: 1.1452 - accuracy: 0.3548  
Epoch 3/50  
8/8 [=====] - 37s 5s/step - loss: 1.1007 - accuracy: 0.3810  
Epoch 4/50  
8/8 [=====] - 37s 5s/step - loss: 1.1034 - accuracy: 0.3750  
Epoch 5/50  
8/8 [=====] - 37s 5s/step - loss: 1.0588 - accuracy: 0.4194  
Epoch 6/50  
8/8 [=====] - 37s 5s/step - loss: 0.9987 - accuracy: 0.4919  
Epoch 7/50  
8/8 [=====] - 37s 5s/step - loss: 0.9502 - accuracy: 0.5060  
Epoch 8/50  
8/8 [=====] - 37s 5s/step - loss: 0.8960 - accuracy: 0.5887  
Epoch 9/50  
8/8 [=====] - 37s 5s/step - loss: 0.8644 - accuracy: 0.5847  
Epoch 10/50  
8/8 [=====] - 37s 5s/step - loss: 0.8516 - accuracy: 0.5847  
Epoch 11/50  
8/8 [=====] - 37s 5s/step - loss: 0.7975 - accuracy: 0.6512  
Epoch 12/50  
8/8 [=====] - 37s 5s/step - loss: 0.7848 - accuracy: 0.6310  
Epoch 13/50  
8/8 [=====] - 37s 5s/step - loss: 0.7357 - accuracy: 0.6532  
Epoch 14/50  
8/8 [=====] - 37s 5s/step - loss: 0.6856 - accuracy: 0.6754  
Epoch 15/50  
8/8 [=====] - 37s 5s/step - loss: 0.6869 - accuracy: 0.6633  
Epoch 16/50  
8/8 [=====] - 37s 5s/step - loss: 0.6399 - accuracy: 0.7117  
Epoch 17/50  
8/8 [=====] - 37s 5s/step - loss: 0.5815 - accuracy: 0.7581  
Epoch 18/50  
8/8 [=====] - 37s 5s/step - loss: 0.5475 - accuracy: 0.7601  
Epoch 19/50  
8/8 [=====] - 37s 5s/step - loss: 0.5016 - accuracy: 0.7823  
Epoch 20/50  
8/8 [=====] - 37s 5s/step - loss: 0.4751 - accuracy: 0.7923  
Epoch 21/50  
8/8 [=====] - 37s 5s/step - loss: 0.3891 - accuracy: 0.8266  
Epoch 22/50  
8/8 [=====] - 37s 5s/step - loss: 0.4010 - accuracy: 0.8367  
Epoch 23/50
```

```
8/8 [=====] - 37s 5s/step - loss: 0.3737 - accuracy: 0.8528
Epoch 24/50
8/8 [=====] - 37s 5s/step - loss: 0.3090 - accuracy: 0.8810
Epoch 25/50
8/8 [=====] - 37s 5s/step - loss: 0.2650 - accuracy: 0.8952
Epoch 26/50
8/8 [=====] - 38s 5s/step - loss: 0.1863 - accuracy: 0.9456
Epoch 27/50
8/8 [=====] - 39s 5s/step - loss: 0.1905 - accuracy: 0.9173
Epoch 28/50
8/8 [=====] - 38s 5s/step - loss: 0.2624 - accuracy: 0.8871
Epoch 29/50
```

```
model.evaluate(test_ds)
```

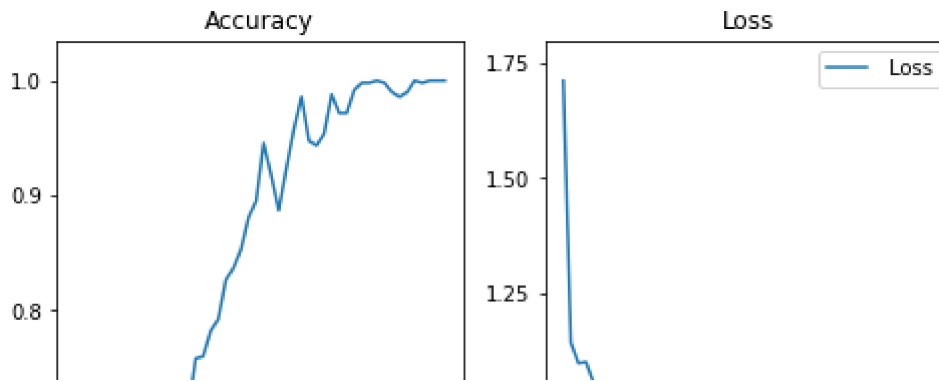
```
1/1 [=====] - 3s 3s/step - loss: 0.3255 - accuracy: 0.9516
[0.325544536113739, 0.9516128897666931]
```

```
acc = history.history['accuracy']
loss = history.history['loss']
```

```
plt.figure(figsize=(8, 8))
plt.subplot(1, 2, 1)
plt.plot(range(EPOCHS), acc, label=' Accuracy')
plt.legend(loc='lower right')
plt.title('Accuracy')
```

```
plt.subplot(1, 2, 2)
plt.plot(range(EPOCHS), loss, label=' Loss')
plt.legend(loc='upper right')
plt.title('Loss')
plt.show()
```





```
image_path = "/gdrive/My Drive/Multinomial animal classification/dogs/dog.5.jpg"
image = preprocessing.image.load_img(image_path)
image_array = preprocessing.image.img_to_array(image)
scaled_img = np.expand_dims(image_array, axis=0)
image
```



```
pred = model.predict(scaled_img)
```

```
output = class_names[np.argmax(pred)]
```

```
output
```

```
'dogs'
```

```
model.save("Multinomial animal classification.h5")
```

